

Discussion L10 (Deep learning representations)

- 1) What is feature engineering? What does learning representation mean, imply? How do DNNs work with data features?
- 2) How is learning representations supposed to support the functionality of DNNs or help in applications? How does it compare to other ANN architectures examined in the course so far?
- 3) What is your intuitive take on the quality of representations? What characteristics a “good” code should have?
- 4) Design a DL system that forms or classifies multi-modal perceptual memories, e.g. images and sounds. What could be a computational basis of multi-modal processing?
- 5) Why would you expect a DNN trained to map the same inputs to different sets of labels (multi-tasking) to learn “better” representations?
- 6) What are relative advantages of sparse and distributed forms of representations? How would you propose to control the level of sparseness?
- 7) Which argument why deep learning offers such fundamental progress in performance on many pattern recognition problems do you find most convincing?
- 8) What is your intuition behind a generative DNN model based approach to learning representations? Is it a promising direction and why? What would you consider a key “selling point” in favour or against this approach?